

▼ Lab 2: Cats vs Dogs

Deadline: Oct 8, 11:59pm

Late Penalty: There is a penalty-free grace period of one hour past the deadline. Any work that is submitted between 1 hour and 24 hours past the deadline will receive a 20% grade deduction. No other late work is accepted. Quercus submission time will be used, not your local computer time. You can submit your labs as many times as you want before the deadline, so please submit often and early.

Marking TA: Tinglin (Francis) Duan

This lab is partially based on an assignment developed by Prof. Jonathan Rose and Harris Chan.

In this lab, you will train a convolutional neural network to classify an image into one of two classes: "cat" or "dog". The code for the neural networks you train will be written for you, and you are not (yet!) expected to understand all provided code. However, by the end of the lab, you should be able to:

1. Understand at a high level the training loop for a machine learning model.
2. Understand the distinction between training, validation, and test data.
3. The concepts of overfitting and underfitting.
4. Investigate how different hyperparameters, such as learning rate and batch size, affect the success of training.
5. Compare an ANN (aka Multi-Layer Perceptron) with a CNN.

What to submit

Submit a PDF file containing all your code, outputs, and write-up from parts 1-5. You can produce a PDF of your Google Colab file by going to **File > Print** and then save as PDF. The Colab instructions has more information.

Do not submit any other files produced by your code.

Include a link to your colab file in your submission.

Please use Google Colab to complete this assignment. If you want to use Jupyter Notebook, please complete the assignment and upload your Jupyter Notebook file to Google Colab for submission.

With Colab, you can export a PDF file using the menu option `File -> Print` and save as PDF file.

Adjust the scaling to ensure that the text is not cutoff at the margins.

▼ Colab Link

Include a link to your colab file here

Colab Link: <https://drive.google.com/file/d/1TMRj0lwWv7w-pMCGXljLQkEreErQYw6a/view?usp=sharing>

```
import numpy as np
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torch.utils.data.sampler import SubsetRandomSampler
import torchvision.transforms as transforms
```

▼ Part 0. Helper Functions

We will be making use of the following helper functions. You will be asked to look at and possibly modify some of these, but you are not expected to understand all of them.

You should look at the function names and read the docstrings. If you are curious, come back and explore the code *after* making some progress on the lab.

```
#####
# Data Loading

def get_relevant_indices(dataset, classes, target_classes):
    """ Return the indices for datapoints in the dataset that belongs to the
        desired target classes, a subset of all possible classes.

    Args:
        dataset: Dataset object
        classes: A list of strings denoting the name of each class
        target_classes: A list of strings denoting the name of desired classes
                       Should be a subset of the 'classes'

    Returns:
        indices: list of indices that have labels corresponding to one of the
                target classes

    """
    indices = []
    for i in range(len(dataset)):
        # Check if the label is in the target classes
        label_index = dataset[i][1] # ex: 3
        label_class = classes[label_index] # ex: 'cat'
        if label_class in target_classes:
            indices.append(i)
    return indices
```

```

def get_data_loader(target_classes, batch_size):
    """ Loads images of cats and dogs, splits the data into training, validation
    and testing datasets. Returns data loaders for the three preprocessed datasets.

    Args:
        target_classes: A list of strings denoting the name of the desired
                        classes. Should be a subset of the argument 'classes'
        batch_size: A int representing the number of samples per batch

    Returns:
        train_loader: iterable training dataset organized according to batch size
        val_loader: iterable validation dataset organized according to batch size
        test_loader: iterable testing dataset organized according to batch size
        classes: A list of strings denoting the name of each class
    """

    classes = ('plane', 'car', 'bird', 'cat',
               'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
    #####
    # The output of torchvision datasets are PILImage images of range [0, 1].
    # We transform them to Tensors of normalized range [-1, 1].
    transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
    # Load CIFAR10 training data
    trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                             download=True, transform=transform)

    # Get the list of indices to sample from
    relevant_indices = get_relevant_indices(trainset, classes, target_classes)

    # Split into train and validation
    np.random.seed(1000) # Fixed numpy random seed for reproducible shuffling
    np.random.shuffle(relevant_indices)
    split = int(len(relevant_indices) * 0.8) #split at 80%

    # split into training and validation indices
    relevant_train_indices, relevant_val_indices = relevant_indices[:split], relevant_indices[split:]
    train_sampler = SubsetRandomSampler(relevant_train_indices)
    train_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                                num_workers=1, sampler=train_sampler)

    val_sampler = SubsetRandomSampler(relevant_val_indices)
    val_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                              num_workers=1, sampler=val_sampler)

    # Load CIFAR10 testing data
    testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                             download=True, transform=transform)

    # Get the list of indices to sample from
    relevant_test_indices = get_relevant_indices(testset, classes, target_classes)
    test_sampler = SubsetRandomSampler(relevant_test_indices)
    test_loader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                              num_workers=1, sampler=test_sampler)

    return train_loader, val_loader, test_loader, classes

```

```

return train_loader, val_loader, test_loader, classes

#####
# Training
def get_model_name(name, batch_size, learning_rate, epoch):
    """ Generate a name for the model consisting of all the hyperparameter values

    Args:
        config: Configuration object containing the hyperparameters
    Returns:
        path: A string with the hyperparameter name and value concatenated
    """
    path = "model_{0}_bs{1}_lr{2}_epoch{3}".format(name,
                                                    batch_size,
                                                    learning_rate,
                                                    epoch)

    return path

def normalize_label(labels):
    """
    Given a tensor containing 2 possible values, normalize this to 0/1

    Args:
        labels: a 1D tensor containing two possible scalar values
    Returns:
        A tensor normalize to 0/1 value
    """
    max_val = torch.max(labels)
    min_val = torch.min(labels)
    norm_labels = (labels - min_val)/(max_val - min_val)
    return norm_labels

def evaluate(net, loader, criterion):
    """ Evaluate the network on the validation set.

    Args:
        net: PyTorch neural network object
        loader: PyTorch data loader for the validation set
        criterion: The loss function
    Returns:
        err: A scalar for the avg classification error over the validation set
        loss: A scalar for the average loss function over the validation set
    """
    total_loss = 0.0
    total_err = 0.0
    total_epoch = 0
    for i, data in enumerate(loader, 0):
        inputs, labels = data
        labels = normalize_label(labels) # Convert labels to 0/1
        outputs = net(inputs)
        loss = criterion(outputs, labels.float())
        corr = (outputs > 0.0).squeeze().long() != labels

```

```

        total_err += int(corr.sum())
        total_loss += loss.item()
        total_epoch += len(labels)
    err = float(total_err) / total_epoch
    loss = float(total_loss) / (i + 1)
    return err, loss

```

```
#####
```

```
# Training Curve
```

```
def plot_training_curve(path):
```

```
    """ Plots the training curve for a model run, given the csv files
    containing the train/validation error/loss.
```

```
    Args:
```

```
        path: The base path of the csv files produced during training
    """
```

```

    import matplotlib.pyplot as plt
    train_err = np.loadtxt("{}_train_err.csv".format(path))
    val_err = np.loadtxt("{}_val_err.csv".format(path))
    train_loss = np.loadtxt("{}_train_loss.csv".format(path))
    val_loss = np.loadtxt("{}_val_loss.csv".format(path))
    plt.title("Train vs Validation Error")
    n = len(train_err) # number of epochs
    plt.plot(range(1,n+1), train_err, label="Train")
    plt.plot(range(1,n+1), val_err, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Error")
    plt.legend(loc='best')
    plt.show()

    plt.title("Train vs Validation Loss")
    plt.plot(range(1,n+1), train_loss, label="Train")
    plt.plot(range(1,n+1), val_loss, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.legend(loc='best')
    plt.show()

```

▼ Part 1. Visualizing the Data [7 pt]

We will make use of some of the CIFAR-10 data set, which consists of colour images of size 32x32 pixels belonging to 10 categories. You can find out more about the dataset at

<https://www.cs.toronto.edu/~kriz/cifar.html>

For this assignment, we will only be using the cat and dog categories. We have included code that automatically downloads the dataset the first time that the main script is run.

```

# This will download the CIFAR-10 dataset to a folder called "data"
# the first time you run this code.
train loader, val loader, test loader, classes = get_data_loader(

```

```
target_classes=["cat", "dog"],
batch_size=1) # One image per batch
```

↳ Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz [00:20<00:00, 61685932.97it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

▼ Part (a) -- 1 pt

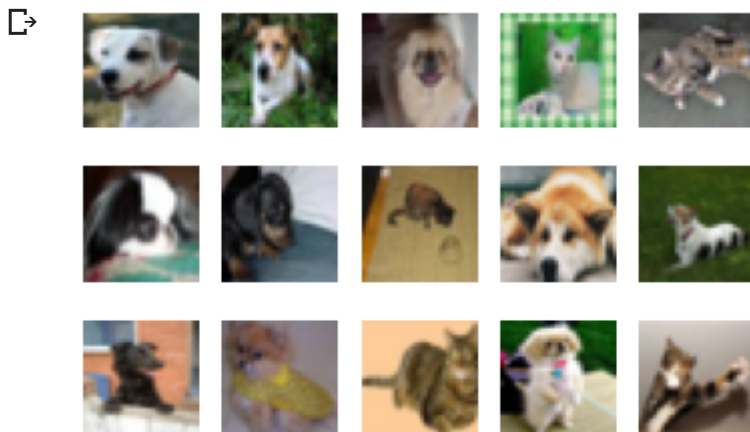
Visualize some of the data by running the code below. Include the visualization in your writeup.

(You don't need to submit anything else.)

```
import matplotlib.pyplot as plt

k = 0
for images, labels in val_loader:
    # since batch_size = 1, there is only 1 image in `images`
    image = images[0]
    # place the colour channel at the end, instead of at the beginning
    img = np.transpose(image, [1,2,0])
    # normalize pixel intensity values to [0, 1]
    img = img / 2 + 0.5
    plt.subplot(3, 5, k+1)
    plt.axis('off')
    plt.imshow(img)

    k += 1
    if k > 14:
        break
```



▼ Part (b) -- 3 pt

How many training examples do we have for the combined `cat` and `dog` classes? What about

```
print(len(train_loader))
print(len(test_loader))
print(len(val_loader))

# There are 8000 training examples and 2000 each of the testing and validation examples
```

```
8000
2000
2000
```

▼ Part (c) -- 3pt

Why do we need a validation set when training our model? What happens if we judge the performance of our models using the training set loss/error instead of the validation set loss/error?

It is important to have a validation set as you want to verify that the model is trained well. There is always the risk of overfitting / underfitting the model, as it may seem like the model works if there is small error or loss. We use the validation set to introduce unseen data to run the model such that we can confirm that it works.

▼ Part 2. Training [15 pt]

We define two neural networks, a `LargeNet` and `SmallNet`. We'll be training the networks in this section.

You won't understand fully what these networks are doing until the next few classes, and that's okay. For this assignment, please focus on learning how to train networks, and how hyperparameters affect training.

```
class LargeNet(nn.Module):
    def __init__(self):
        super(LargeNet, self).__init__()
        self.name = "large"
        self.conv1 = nn.Conv2d(3, 5, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(5, 10, 5)
        self.fc1 = nn.Linear(10 * 5 * 5, 32)
        self.fc2 = nn.Linear(32, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 10 * 5 * 5)
        x = self.fc1(x)
        x = self.fc2(x)
```

```

x = F.relu(self.fc1(x))
x = self.fc2(x)
x = x.squeeze(1) # Flatten to [batch_size]
return x

```

```

class SmallNet(nn.Module):
    def __init__(self):
        super(SmallNet, self).__init__()
        self.name = "small"
        self.conv = nn.Conv2d(3, 5, 3)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc = nn.Linear(5 * 7 * 7, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv(x)))
        x = self.pool(x)
        x = x.view(-1, 5 * 7 * 7)
        x = self.fc(x)
        x = x.squeeze(1) # Flatten to [batch_size]
        return x

small_net = SmallNet()
large_net = LargeNet()

```

▼ Part (a) -- 2pt

The methods `small_net.parameters()` and `large_net.parameters()` produces an iterator of all the trainable parameters of the network. These parameters are torch tensors containing many scalar values.

We haven't learned how the parameters in these high-dimensional tensors will be used, but we should be able to count the number of parameters. Measuring the number of parameters in a network is one way of measuring the "size" of a network.

What is the total number of parameters in `small_net` and in `large_net`? (Hint: how many numbers are in each tensor?)

```

accum = 1
for param in small_net.parameters():
    for i in range(0, len(param.shape)):
        accum = accum*param.shape[i]

print("There are " + str(accum) + " parameters in the small_net")

accum = 1
for param in large_net.parameters():
    for i in range(0, len(param.shape)):
        accum = accum*param.shape[i]

```



```
print("There are " + str(accum) + " parameters in the small_net")
```

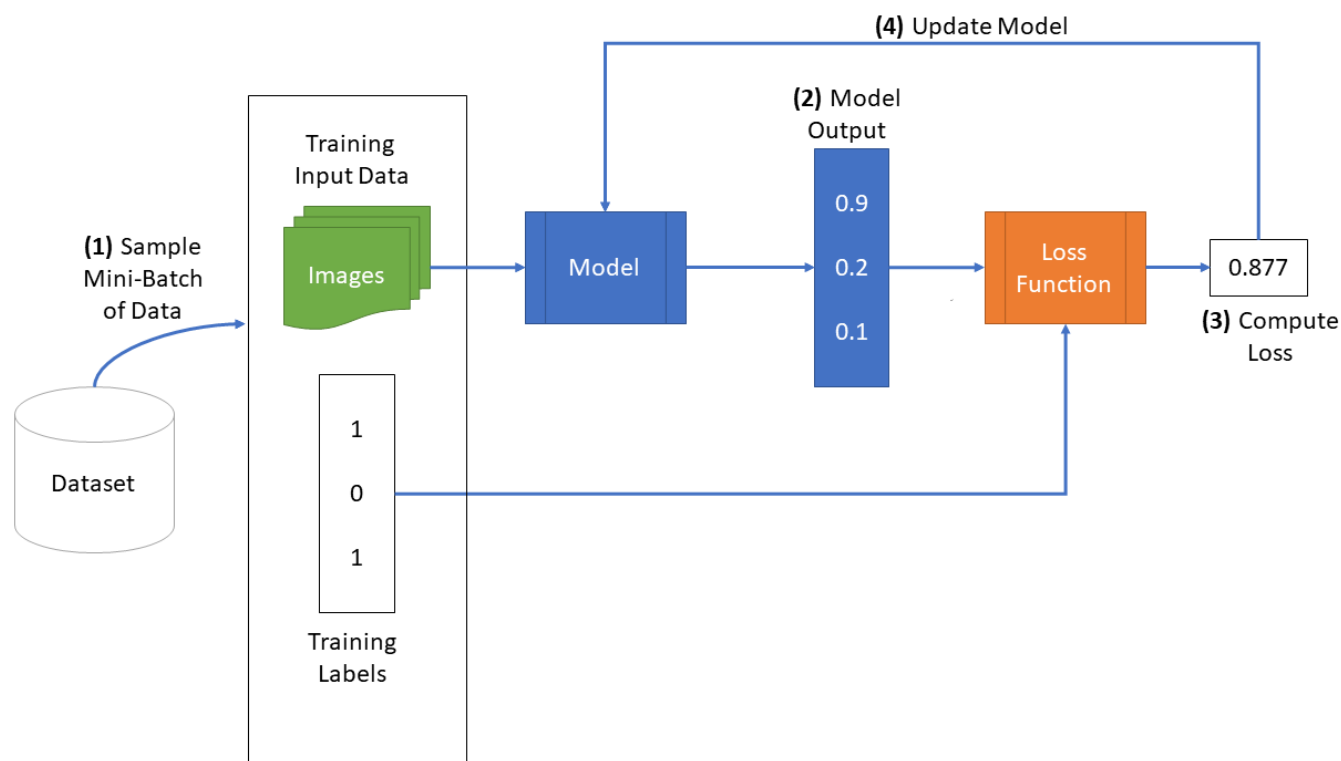
```

[ ] There are 165375 parameters in the small_net
    There are 192000000000000 parameters in the small_net

```

▼ The function train_net

The function `train_net` below takes an untrained neural network (like `small_net` and `large_net`) and several other parameters. You should be able to understand how this function works. The figure below shows the high level training loop for a machine learning model:



```

def train_net(net, batch_size=64, learning_rate=0.01, num_epochs=30):
    #####
    # Train a classifier on cats vs dogs
    target_classes = ["cat", "dog"]
    #####
    # Fixed PyTorch random seed for reproducible result
    torch.manual_seed(1000)
    #####
    # Obtain the PyTorch data loader objects to load batches of the datasets
    train_loader, val_loader, test_loader, classes = get_data_loader(
        target_classes, batch_size)
    #####
    # Define the Loss function and optimizer
    # The loss function will be Binary Cross Entropy (BCE). In this case we
    "

```

```

# Will use the BCEWithLogitsLoss which takes unnormalized output from
# the neural network and scalar label.
# Optimizer will be SGD with Momentum.
criterion = nn.BCEWithLogitsLoss()
optimizer = optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
#####
# Set up some numpy arrays to store the training/test loss/erruracy
train_err = np.zeros(num_epochs)
train_loss = np.zeros(num_epochs)
val_err = np.zeros(num_epochs)
val_loss = np.zeros(num_epochs)
#####
# Train the network
# Loop over the data iterator and sample a new batch of training data
# Get the output from the network, and optimize our loss function.
start_time = time.time()
for epoch in range(num_epochs): # loop over the dataset multiple times
    total_train_loss = 0.0
    total_train_err = 0.0
    total_epoch = 0
    for i, data in enumerate(train_loader, 0):
        # Get the inputs
        inputs, labels = data
        labels = normalize_label(labels) # Convert labels to 0/1
        # Zero the parameter gradients
        optimizer.zero_grad()
        # Forward pass, backward pass, and optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels.float())
        loss.backward()
        optimizer.step()
        # Calculate the statistics
        corr = (outputs > 0.0).squeeze().long() != labels
        total_train_err += int(corr.sum())
        total_train_loss += loss.item()
        total_epoch += len(labels)
    train_err[epoch] = float(total_train_err) / total_epoch
    train_loss[epoch] = float(total_train_loss) / (i+1)
    val_err[epoch], val_loss[epoch] = evaluate(net, val_loader, criterion)
    print("Epoch {}: Train err: {}, Train loss: {} |"+
          "Validation err: {}, Validation loss: {}".format(
                epoch + 1,
                train_err[epoch],
                train_loss[epoch],
                val_err[epoch],
                val_loss[epoch]))
    # Save the current model (checkpoint) to a file
    model_path = get_model_name(net.name, batch_size, learning_rate, epoch)
    torch.save(net.state_dict(), model_path)
print('Finished Training')
end_time = time.time()
elapsed_time = end_time - start_time

```

```
print("Total time elapsed: {:.2f} seconds".format(elapsed_time))
# Write the train/test loss/err into CSV file for plotting later
epochs = np.arange(1, num_epochs + 1)
np.savetxt("{}_train_err.csv".format(model_path), train_err)
np.savetxt("{}_train_loss.csv".format(model_path), train_loss)
np.savetxt("{}_val_err.csv".format(model_path), val_err)
np.savetxt("{}_val_loss.csv".format(model_path), val_loss)
```

▼ Part (b) -- 1pt

The parameters to the function `train_net` are hyperparameters of our neural network. We made these hyperparameters easy to modify so that we can tune them later on.

What are the default values of the parameters `batch_size`, `learning_rate`, and `num_epochs`?

```
# batch_size = 64
# learning_rate = 0.01
# num_epochs = 30
```

▼ Part (c) -- 3 pt

What files are written to disk when we call `train_net` with `small_net`, and train for 5 epochs?

Provide a list of all the files written to disk, and what information the files contain.

```
train_net(small_net, 64, 0.01, 5)
```

```
☞ Files already downloaded and verified
Files already downloaded and verified
Epoch 1: Train err: 0.418125, Train loss: 0.6704066848754883 | Validation err: 0.33425
Epoch 2: Train err: 0.364875, Train loss: 0.6426583008766175 | Validation err: 0.33425
Epoch 3: Train err: 0.353375, Train loss: 0.6322054648399353 | Validation err: 0.33425
Epoch 4: Train err: 0.342375, Train loss: 0.6217929873466491 | Validation err: 0.33425
Epoch 5: Train err: 0.33425, Train loss: 0.6123823966979981 | Validation err: 0.33425
Finished Training
Total time elapsed: 18.50 seconds
```

The following files written to disk:

```
model_small_bs64_lr0.01_epoch0
model_small_bs64_lr0.01_epoch1
model_small_bs64_lr0.01_epoch2
model_small_bs64_lr0.01_epoch3
model_small_bs64_lr0.01_epoch4
model_small_bs64_lr0.01_epoch4_train_err
model_small_bs64_lr0.01_epoch_train_loss
```

```
model_small_bs64_lr0.01_epoch_val_loss
```

```
model_small_bs64_lr0.01_epoch_val_err
```

▼ Part (d) -- 2pt

Train both `small_net` and `large_net` using the function `train_net` and its default parameters. The function will write many files to disk, including a model checkpoint (saved values of model weights) at the end of each epoch.

If you are using Google Colab, you will need to mount Google Drive so that the files generated by `train_net` gets saved. We will be using these files in part (d). (See the Google Colab tutorial for more information about this.)

Report the total time elapsed when training each network. Which network took longer to train? Why?

```
# Since the function writes files to disk, you will need to mount
# your Google Drive. If you are working on the lab locally, you
# can comment out this code.
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
↳ Mounted at /content/gdrive
```

```
train_net(small_net, 64, 0.01, 30)
## 111.05 seconds to train
```

```
train_net(large_net, 64, 0.01, 30)
## 124.64 seconds to train
```

```
# The large_net took longer as there are more parameters to train
```

```
↳
```

Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.32825, Train loss: 0.6051286149024964 |Validation err: 0.3
Epoch 2: Train err: 0.31775, Train loss: 0.5988569338321685 |Validation err: 0.3
Epoch 3: Train err: 0.3205, Train loss: 0.5974532868862152 |Validation err: 0.33
Epoch 4: Train err: 0.320125, Train loss: 0.5951218860149383 |Validation err: 0.
Epoch 5: Train err: 0.31175, Train loss: 0.5874869222640992 |Validation err: 0.3
Epoch 6: Train err: 0.305375, Train loss: 0.5835547342300414 |Validation err: 0.
Epoch 7: Train err: 0.307, Train loss: 0.5832296054363251 |Validation err: 0.329
Epoch 8: Train err: 0.306125, Train loss: 0.5778174822330475 |Validation err: 0.
Epoch 9: Train err: 0.30175, Train loss: 0.5817056109905243 |Validation err: 0.3
Epoch 10: Train err: 0.298875, Train loss: 0.5740554056167603 |Validation err: 0
Epoch 11: Train err: 0.295625, Train loss: 0.5715680284500122 |Validation err: 0
Epoch 12: Train err: 0.297125, Train loss: 0.5660525941848755 |Validation err: 0
Epoch 13: Train err: 0.29225, Train loss: 0.5667166328430175 |Validation err: 0.
Epoch 14: Train err: 0.293125, Train loss: 0.5630647814273835 |Validation err: 0
Epoch 15: Train err: 0.291625, Train loss: 0.5624129445552826 |Validation err: 0
Epoch 16: Train err: 0.294875, Train loss: 0.5652327179908753 |Validation err: 0
Epoch 17: Train err: 0.293625, Train loss: 0.5638533592224121 |Validation err: 0
Epoch 18: Train err: 0.285125, Train loss: 0.555136979341507 |Validation err: 0.
Epoch 19: Train err: 0.28125, Train loss: 0.5518331458568573 |Validation err: 0.
Epoch 20: Train err: 0.287375, Train loss: 0.5528540678024292 |Validation err: 0
Epoch 21: Train err: 0.277625, Train loss: 0.5527849159240723 |Validation err: 0
Epoch 22: Train err: 0.281875, Train loss: 0.5523663868904114 |Validation err: 0
Epoch 23: Train err: 0.276875, Train loss: 0.5489318742752075 |Validation err: 0
Epoch 24: Train err: 0.275875, Train loss: 0.5477618153095245 |Validation err: 0
Epoch 25: Train err: 0.274375, Train loss: 0.5441374328136445 |Validation err: 0
Epoch 26: Train err: 0.273625, Train loss: 0.542892305135727 |Validation err: 0.
Epoch 27: Train err: 0.278, Train loss: 0.5421670744419098 |Validation err: 0.29
Epoch 28: Train err: 0.278375, Train loss: 0.5446225323677063 |Validation err: 0
Epoch 29: Train err: 0.275625, Train loss: 0.542576758146286 |Validation err: 0.
Epoch 30: Train err: 0.278, Train loss: 0.5443889632225036 |Validation err: 0.29
```

Finished Training

Total time elapsed: 108.04 seconds

Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.471625, Train loss: 0.6914801979064942 |Validation err: 0.
Epoch 2: Train err: 0.435125, Train loss: 0.6820778837203979 |Validation err: 0.
Epoch 3: Train err: 0.406, Train loss: 0.6679732670783997 |Validation err: 0.373
Epoch 4: Train err: 0.36675, Train loss: 0.6430274744033814 |Validation err: 0.3
Epoch 5: Train err: 0.350125, Train loss: 0.6276090312004089 |Validation err: 0.
Epoch 6: Train err: 0.330625, Train loss: 0.6105496971607208 |Validation err: 0.
Epoch 7: Train err: 0.322125, Train loss: 0.6005237801074982 |Validation err: 0.
Epoch 8: Train err: 0.30875, Train loss: 0.5852175650596618 |Validation err: 0.3
Epoch 9: Train err: 0.309125, Train loss: 0.5838861713409423 |Validation err: 0.
Epoch 10: Train err: 0.294875, Train loss: 0.5673127889633178 |Validation err: 0
Epoch 11: Train err: 0.284, Train loss: 0.5562010607719421 |Validation err: 0.32
Epoch 12: Train err: 0.27275, Train loss: 0.5404626941680908 |Validation err: 0.
Epoch 13: Train err: 0.27075, Train loss: 0.5338843922615051 |Validation err: 0.
Epoch 14: Train err: 0.262125, Train loss: 0.5188266251087189 |Validation err: 0
Epoch 15: Train err: 0.256125, Train loss: 0.5122005236148834 |Validation err: 0
Epoch 16: Train err: 0.247, Train loss: 0.5076880040168762 |Validation err: 0.29
Epoch 17: Train err: 0.2425, Train loss: 0.49078083753585816 |Validation err: 0.
```

▼ Part (e) - 2pt

Use the function `plot_training_curve` to display the trajectory of the training/validation error and the training/validation loss. You will need to use the function `get_model_name` to generate the argument to the `plot_training_curve` function.

Do this for both the small network and the large network. Include both plots in your writeup.

```
Epoch 28: Train err: 0.16225, Train loss: 0.3586689592599869 | Validation err: 0.1
```

```
model_path = get_model_name("small", batch_size=64, learning_rate=0.01, epoch=29)
plot_training_curve(model_path)
```

```
model_path = get_model_name("large", batch_size=64, learning_rate=0.01, epoch=29)
plot_training_curve(model_path)
```





▼ Part (f) - 5pt

Describe what you notice about the training curve. How do the curves differ for `small_net` and `large_net`? Identify any occurrences of underfitting and overfitting.

|  |

For the `small_net`, the training curves for both the error and the loss are lower than the validation curves, meaning less error and loss. In both cases, the training curves also trend downward as the epoch increase to smaller values while the validation curves plateau.


For the `large_net`, the training curves match up very well with the validation curves while the epoch number is low. As the epochs increase, the two curves tend to diverge, where the training curve continues on a downward linear trend and the validation curve trend in the opposite direction.

As the training error decreases and the validation error does not, these cases are considered to be overfitting. This can be seen in the areas where Epoch > 15. In the areas where Epoch < 10, the model is underfitting as the training error continues to decrease even though the complexity is increasing with the epochs.

--- |  |

▼ Part 3. Optimization Parameters [12 pt]

For this section, we will work with `large_net` only.

|  |  |  |

▼ Part (a) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.001`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *lowering* the learning rate.

0.40 ↓

|

```
# Note: When we re-construct the model, we start the training
# with *random weights*. If we omit this code, the values of
# the weights will still be the previously trained values.
large_net = LargeNet()

train_net(large_net, 64, 0.001, 30)
model_path = get_model_name("large", batch_size=64, learning_rate=0.001, epoch=29)
plot_training_curve(model_path)
```



Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.47625, Train loss: 0.6928360004425049 | Validation err: 0.4
Epoch 2: Train err: 0.448625, Train loss: 0.6922589735984802 | Validation err: 0.
Epoch 3: Train err: 0.43575, Train loss: 0.6916067390441895 | Validation err: 0.4
Epoch 4: Train err: 0.430125, Train loss: 0.6908613820075988 | Validation err: 0.
Epoch 5: Train err: 0.434125, Train loss: 0.6899198398590088 | Validation err: 0.
Epoch 6: Train err: 0.435875, Train loss: 0.6887419700622559 | Validation err: 0.
Epoch 7: Train err: 0.436625, Train loss: 0.6873781762123108 | Validation err: 0.
Epoch 8: Train err: 0.43725, Train loss: 0.6859267873764038 | Validation err: 0.4
Epoch 9: Train err: 0.424375, Train loss: 0.6844043645858765 | Validation err: 0.
Epoch 10: Train err: 0.42425, Train loss: 0.6828487544059754 | Validation err: 0.
Epoch 11: Train err: 0.42525, Train loss: 0.681236349105835 | Validation err: 0.4
Epoch 12: Train err: 0.419875, Train loss: 0.6796339492797852 | Validation err: 0
Epoch 13: Train err: 0.41475, Train loss: 0.6777917776107788 | Validation err: 0.
Epoch 14: Train err: 0.412375, Train loss: 0.6761110095977784 | Validation err: 0
Epoch 15: Train err: 0.40925, Train loss: 0.674471221446991 | Validation err: 0.4
Epoch 16: Train err: 0.4065, Train loss: 0.6727405357360839 | Validation err: 0.4
Epoch 17: Train err: 0.4015, Train loss: 0.6713058156967163 | Validation err: 0.4
Epoch 18: Train err: 0.39925, Train loss: 0.6696760263442993 | Validation err: 0.
Epoch 19: Train err: 0.400875, Train loss: 0.6679068713188171 | Validation err: 0
Epoch 20: Train err: 0.392375, Train loss: 0.6657861313819885 | Validation err: 0
Epoch 21: Train err: 0.3895, Train loss: 0.6646247744560242 | Validation err: 0.3
Epoch 22: Train err: 0.38875, Train loss: 0.6623678812980652 | Validation err: 0.
Epoch 23: Train err: 0.384, Train loss: 0.6601416163444519 | Validation err: 0.39
Epoch 24: Train err: 0.382625, Train loss: 0.6583918170928955 | Validation err: 0
```

The model takes slightly longer to train than before. We notice with the lower learning rate that the training and validation curves are much more similar. The overall error continues to decrease with the increasing epochs, but even with the lower learning rate of 0.001, the error was still higher.

Essentially, having a lower learning rate does not increase efficiency, as the run time is greater and the error is actually greater, so this could be considered to be ineffective.

| 

▼ Part (b) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.1`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the learning rate.

0.38 ↓

 |

```
large_net = LargeNet()
train_net(large_net, 64, 0.1, 30)
model_path = get_model_name("large", batch_size=64, learning_rate=0.1, epoch=29)
plot_training_curve(model_path)
```

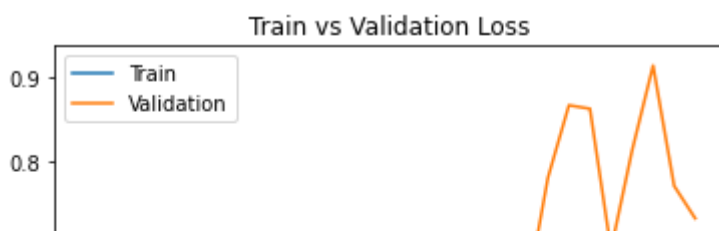
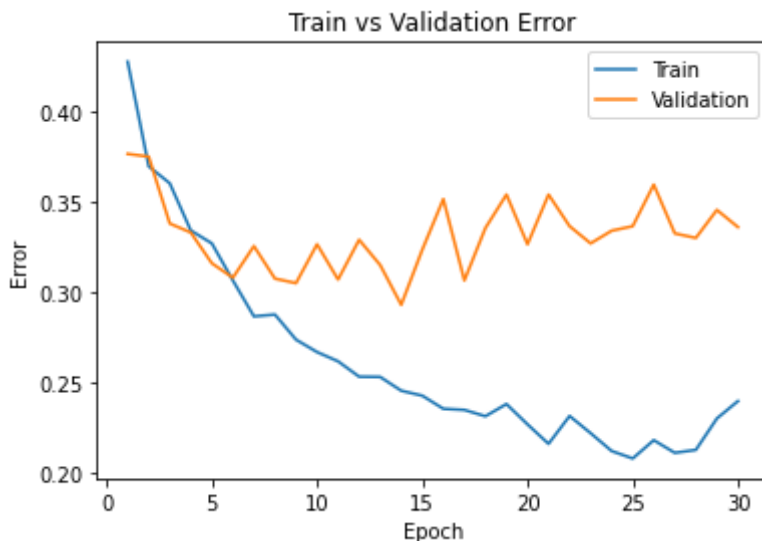


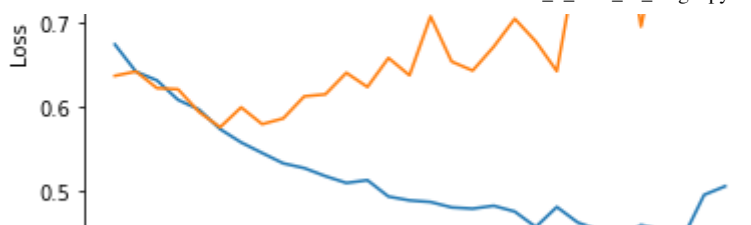
Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.4275, Train loss: 0.6742977547645569 |Validation err: 0.370
Epoch 2: Train err: 0.369625, Train loss: 0.6419942178726197 |Validation err: 0.3
Epoch 3: Train err: 0.36025, Train loss: 0.6315926456451416 |Validation err: 0.3
Epoch 4: Train err: 0.334125, Train loss: 0.6084355125427247 |Validation err: 0.3
Epoch 5: Train err: 0.327, Train loss: 0.5968132178783416 |Validation err: 0.316
Epoch 6: Train err: 0.3065, Train loss: 0.5736211636066437 |Validation err: 0.30
Epoch 7: Train err: 0.286625, Train loss: 0.557759182214737 |Validation err: 0.3
Epoch 8: Train err: 0.287625, Train loss: 0.5451903882026672 |Validation err: 0.3
Epoch 9: Train err: 0.27375, Train loss: 0.5328604147434235 |Validation err: 0.3
Epoch 10: Train err: 0.266875, Train loss: 0.5272168016433716 |Validation err: 0.3
Epoch 11: Train err: 0.26175, Train loss: 0.5176891093254089 |Validation err: 0.3
Epoch 12: Train err: 0.25325, Train loss: 0.5095298163890839 |Validation err: 0.3
Epoch 13: Train err: 0.253125, Train loss: 0.5128449490070343 |Validation err: 0.3
Epoch 14: Train err: 0.2455, Train loss: 0.49335162925720216 |Validation err: 0.3
Epoch 15: Train err: 0.24275, Train loss: 0.48878711795806884 |Validation err: 0.3
Epoch 16: Train err: 0.2355, Train loss: 0.48690388107299803 |Validation err: 0.3
Epoch 17: Train err: 0.234875, Train loss: 0.48044825553894044 |Validation err: 0.3
Epoch 18: Train err: 0.231375, Train loss: 0.4789571635723114 |Validation err: 0.3
Epoch 19: Train err: 0.238125, Train loss: 0.4823964536190033 |Validation err: 0.3
Epoch 20: Train err: 0.226875, Train loss: 0.4755194764137268 |Validation err: 0.3
Epoch 21: Train err: 0.216125, Train loss: 0.45714300179481504 |Validation err: 0.3
Epoch 22: Train err: 0.2315, Train loss: 0.48096809661388396 |Validation err: 0.3
Epoch 23: Train err: 0.222, Train loss: 0.46243434739112854 |Validation err: 0.3
Epoch 24: Train err: 0.212125, Train loss: 0.4531472542285919 |Validation err: 0.3
Epoch 25: Train err: 0.208, Train loss: 0.4432792100906372 |Validation err: 0.33
Epoch 26: Train err: 0.218125, Train loss: 0.45924149703979494 |Validation err: 0.3
Epoch 27: Train err: 0.211125, Train loss: 0.4533423124551773 |Validation err: 0.3
Epoch 28: Train err: 0.21275, Train loss: 0.44430617237091063 |Validation err: 0.3
Epoch 29: Train err: 0.230125, Train loss: 0.4953969626426697 |Validation err: 0.3
Epoch 30: Train err: 0.23975, Train loss: 0.5056364586353302 |Validation err: 0.3
Finished Training
```

Total time elapsed: 122.79 seconds





The model takes nearly the same amount of time to run with the higher learning rate. We see a trend of decreasing validation loss and error up until about epoch = 5, when it then begins to reverse its trend to an upward direction diverging from the training error and loss.

We also see that the change in error and loss as the epochs increase is much more significant with this higher learning rate of 0.1, compared to the learning rate of 0.01.

▼ Part (c) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=512`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the batch size.

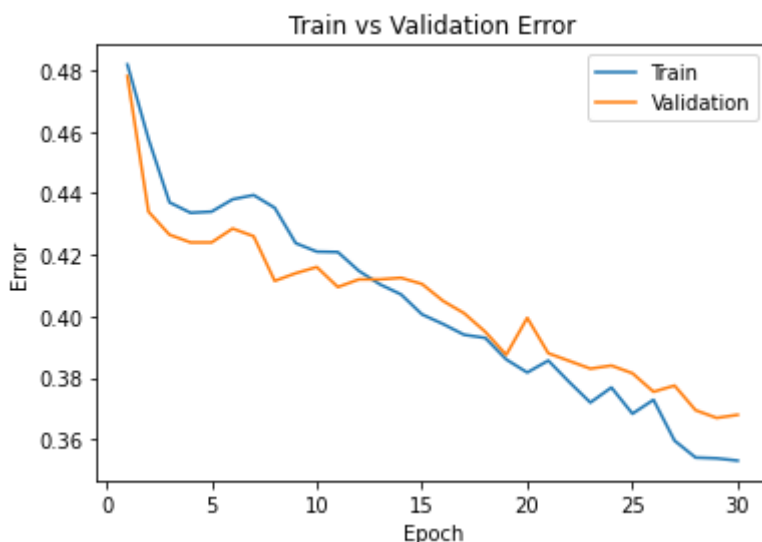
```
large_net = LargeNet()
train_net(large_net, 512, 0.01, 30)
model_path = get_model_name("large", batch_size=512, learning_rate=0.01, epoch=29)
plot_training_curve(model_path)
```



```

Epoch 7: Train err: 0.439375, Train loss: 0.6866871640086174 | Validation err: 0.439375
Epoch 8: Train err: 0.43525, Train loss: 0.6849770732223988 | Validation err: 0.43525
Epoch 9: Train err: 0.42375, Train loss: 0.6832009218633175 | Validation err: 0.42375
Epoch 10: Train err: 0.421, Train loss: 0.6811090856790543 | Validation err: 0.421
Epoch 11: Train err: 0.420875, Train loss: 0.6794028840959072 | Validation err: 0.420875
Epoch 12: Train err: 0.41475, Train loss: 0.676804918795824 | Validation err: 0.41475
Epoch 13: Train err: 0.410375, Train loss: 0.6749706864356995 | Validation err: 0.410375
Epoch 14: Train err: 0.407125, Train loss: 0.6730913147330284 | Validation err: 0.407125
Epoch 15: Train err: 0.400625, Train loss: 0.6706849597394466 | Validation err: 0.400625
Epoch 16: Train err: 0.3975, Train loss: 0.6691824123263359 | Validation err: 0.3975
Epoch 17: Train err: 0.394, Train loss: 0.6675699539482594 | Validation err: 0.394
Epoch 18: Train err: 0.393, Train loss: 0.6647988669574261 | Validation err: 0.393
Epoch 19: Train err: 0.386, Train loss: 0.6627439819276333 | Validation err: 0.386
Epoch 20: Train err: 0.38175, Train loss: 0.6596207581460476 | Validation err: 0.38175
Epoch 21: Train err: 0.385625, Train loss: 0.6584848575294018 | Validation err: 0.385625
Epoch 22: Train err: 0.378625, Train loss: 0.6551308929920197 | Validation err: 0.378625
Epoch 23: Train err: 0.372, Train loss: 0.650898166000843 | Validation err: 0.372
Epoch 24: Train err: 0.376875, Train loss: 0.6488463133573532 | Validation err: 0.376875
Epoch 25: Train err: 0.368375, Train loss: 0.6446415856480598 | Validation err: 0.368375
Epoch 26: Train err: 0.372875, Train loss: 0.6429142691195011 | Validation err: 0.372875
Epoch 27: Train err: 0.359625, Train loss: 0.6373121663928032 | Validation err: 0.359625
Epoch 28: Train err: 0.354125, Train loss: 0.6338207013905048 | Validation err: 0.354125
Epoch 29: Train err: 0.353875, Train loss: 0.6311671584844589 | Validation err: 0.353875
Epoch 30: Train err: 0.353125, Train loss: 0.6283805817365646 | Validation err: 0.353125
Finished Training
Total time elapsed: 107.30 seconds

```



Now, by increasing the batch size to 512, the model takes significantly less time to train, it being approximately 20% more efficient. The training and validation curves for both error and loss are very similar and they trend consistently negative and appear to be quite smooth.

0.67 | |

▼ Part (d) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=16`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *decreasing* the batch size.

```
large_net = LargeNet()  
train_net(large_net, 16, 0.01, 30)  
model_path = get_model_name("large", batch_size=16, learning_rate=0.01, epoch=29)  
plot_training_curve(model_path)
```



```

Files already downloaded and verified
Files already downloaded and verified
Epoch 1: Train err: 0.432625, Train loss: 0.6775410465598106 |Validation err: 0.1
Epoch 2: Train err: 0.367875, Train loss: 0.6389183864593506 |Validation err: 0.1
Epoch 3: Train err: 0.340375, Train loss: 0.6114947981238366 |Validation err: 0.1
Epoch 4: Train err: 0.31175, Train loss: 0.5834803532958031 |Validation err: 0.3
Epoch 5: Train err: 0.301125, Train loss: 0.5638278910517692 |Validation err: 0.1
Epoch 6: Train err: 0.28475, Train loss: 0.5464649626612663 |Validation err: 0.3
Epoch 7: Train err: 0.270625, Train loss: 0.5266695675551891 |Validation err: 0.1
Epoch 8: Train err: 0.254, Train loss: 0.5092709631323814 |Validation err: 0.322
Epoch 9: Train err: 0.245625, Train loss: 0.4968788513541222 |Validation err: 0.1
Epoch 10: Train err: 0.231875, Train loss: 0.4726636277139187 |Validation err: 0.1
Epoch 11: Train err: 0.2195, Train loss: 0.45416154369711875 |Validation err: 0.1
Epoch 12: Train err: 0.2135, Train loss: 0.43952479952573775 |Validation err: 0.1
Epoch 13: Train err: 0.198875, Train loss: 0.41390844713151453 |Validation err: 0.1
Epoch 14: Train err: 0.185375, Train loss: 0.3907764309346676 |Validation err: 0.1
Epoch 15: Train err: 0.174625, Train loss: 0.37488927601277827 |Validation err: 0.1
Epoch 16: Train err: 0.163125, Train loss: 0.3653211060985923 |Validation err: 0.1

```

With a batch size of 16 this time, the validation curves diverged from the training curves by the time the epoch was about 5. This is not necessarily a good sign as it could signal that the model was overfitting as the training error continues to decrease while the validation error is about 3 times as large at 30 epochs.

```

Epoch 24: Train err: 0.164125, Train loss: 0.3478218918228923 |Validation err: 0.1

```

▼ Part 4. Hyperparameter Search [6 pt]

Part (a) - 2pt

Based on the plots from above, choose another set of values for the hyperparameters (network, batch_size, learning_rate) that you think would help you improve the validation accuracy. Justify your choice.

0.40 | \ | Validation |

Based on the plots above, I would choose a moderate batch size of 128 with a lower learning rate of 0.005. Lowering the batch size earlier didn't really help and setting it extremely high wasn't amazing either, so I'd like to try a moderate level. Additionally, a lower learning rate would move things in a smaller step which could improve a accuracy

| |

▼ Part (b) - 1pt

Train the model with the hyperparameters you chose in part(a), and include the training curve.

```

from

```

```

large_net = LargeNet()

train_net(large_net, batch_size=128, learning_rate=0.005, num_epochs=30)
model_path = get_model_name("large", batch_size=128, learning_rate=0.005, epoch=29)
plot_training_curve(model_path)

```

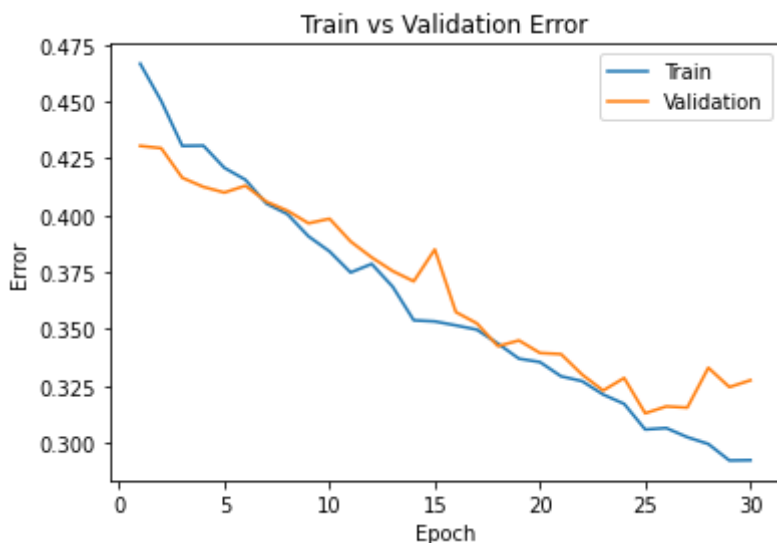


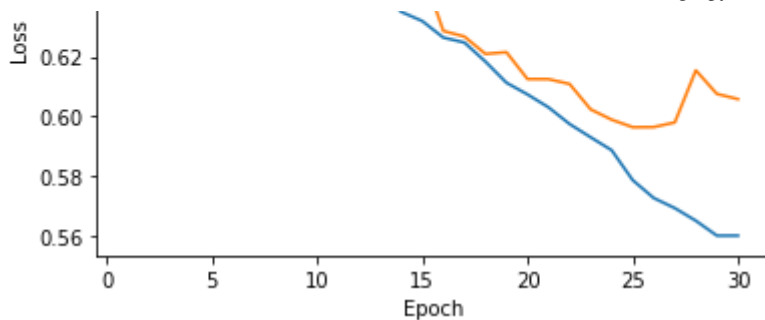
Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.466625, Train loss: 0.6925613265188914 |Validation err: 0.466625, Train loss: 0.6925613265188914
Epoch 2: Train err: 0.45025, Train loss: 0.6910346150398254 |Validation err: 0.45025, Train loss: 0.6910346150398254
Epoch 3: Train err: 0.4305, Train loss: 0.6885887715551589 |Validation err: 0.4305, Train loss: 0.6885887715551589
Epoch 4: Train err: 0.430625, Train loss: 0.6850015464283171 |Validation err: 0.430625, Train loss: 0.6850015464283171
Epoch 5: Train err: 0.420875, Train loss: 0.6813878568391951 |Validation err: 0.420875, Train loss: 0.6813878568391951
Epoch 6: Train err: 0.415625, Train loss: 0.6772933602333069 |Validation err: 0.415625, Train loss: 0.6772933602333069
Epoch 7: Train err: 0.405125, Train loss: 0.6732233772202144 |Validation err: 0.405125, Train loss: 0.6732233772202144
Epoch 8: Train err: 0.4005, Train loss: 0.6694089571634928 |Validation err: 0.4005, Train loss: 0.6694089571634928
Epoch 9: Train err: 0.39075, Train loss: 0.665704798130762 |Validation err: 0.39075, Train loss: 0.665704798130762
Epoch 10: Train err: 0.384125, Train loss: 0.6609769766292874 |Validation err: 0.384125, Train loss: 0.6609769766292874
Epoch 11: Train err: 0.374875, Train loss: 0.6547541533197675 |Validation err: 0.374875, Train loss: 0.6547541533197675
Epoch 12: Train err: 0.378625, Train loss: 0.6503568007832482 |Validation err: 0.378625, Train loss: 0.6503568007832482
Epoch 13: Train err: 0.368625, Train loss: 0.6430440866757953 |Validation err: 0.368625, Train loss: 0.6430440866757953
Epoch 14: Train err: 0.353875, Train loss: 0.6349754617327735 |Validation err: 0.353875, Train loss: 0.6349754617327735
Epoch 15: Train err: 0.353375, Train loss: 0.631892790870061 |Validation err: 0.353375, Train loss: 0.631892790870061
Epoch 16: Train err: 0.351625, Train loss: 0.6263627977598281 |Validation err: 0.351625, Train loss: 0.6263627977598281
Epoch 17: Train err: 0.34975, Train loss: 0.6247141408541846 |Validation err: 0.34975, Train loss: 0.6247141408541846
Epoch 18: Train err: 0.343625, Train loss: 0.6183376870458088 |Validation err: 0.343625, Train loss: 0.6183376870458088
Epoch 19: Train err: 0.337, Train loss: 0.6112656962303888 |Validation err: 0.337, Train loss: 0.6112656962303888
Epoch 20: Train err: 0.3355, Train loss: 0.6073562540705242 |Validation err: 0.3355, Train loss: 0.6073562540705242
Epoch 21: Train err: 0.32925, Train loss: 0.6030320581935701 |Validation err: 0.32925, Train loss: 0.6030320581935701
Epoch 22: Train err: 0.327125, Train loss: 0.5974628300893874 |Validation err: 0.327125, Train loss: 0.5974628300893874
Epoch 23: Train err: 0.32125, Train loss: 0.5930387093907311 |Validation err: 0.32125, Train loss: 0.5930387093907311
Epoch 24: Train err: 0.317125, Train loss: 0.5886963378815424 |Validation err: 0.317125, Train loss: 0.5886963378815424
Epoch 25: Train err: 0.305875, Train loss: 0.5786937678617144 |Validation err: 0.305875, Train loss: 0.5786937678617144
Epoch 26: Train err: 0.3065, Train loss: 0.5727116952812861 |Validation err: 0.3065, Train loss: 0.5727116952812861
Epoch 27: Train err: 0.3025, Train loss: 0.5693432674521491 |Validation err: 0.3025, Train loss: 0.5693432674521491
Epoch 28: Train err: 0.2995, Train loss: 0.5651203800761511 |Validation err: 0.2995, Train loss: 0.5651203800761511
Epoch 29: Train err: 0.29225, Train loss: 0.5601372870187911 |Validation err: 0.29225, Train loss: 0.5601372870187911
Epoch 30: Train err: 0.292375, Train loss: 0.5601449537844885 |Validation err: 0.292375, Train loss: 0.5601449537844885
Finished Training
```

Total time elapsed: 109.62 seconds





▼ Part (c) - 2pt

Based on your result from Part(a), suggest another set of hyperparameter values to try. Justify your choice.

Based on the results above, this time I'd like to try lowering the learning rate a little bit more to see if there is any more improvement. This time I will try 0.0001. The validation curve looks good until about the 25 epoch mark where it slightly diverges and plateaus.

▼ Part (d) - 1pt

Train the model with the hyperparameters you chose in part(c), and include the training curve.

```
large_net = LargeNet()

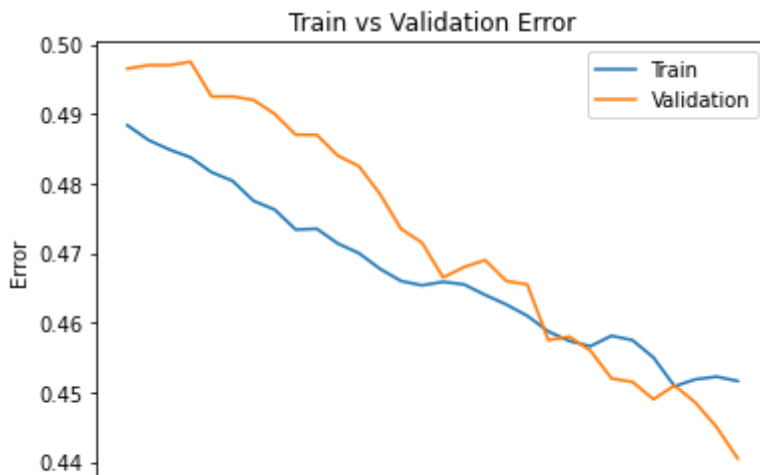
train_net(large_net, batch_size=128, learning_rate=0.0001, num_epochs=30)
model_path = get_model_name("large", batch_size=128, learning_rate=0.0001, epoch=29)
plot_training_curve(model_path)
```



Files already downloaded and verified

Files already downloaded and verified

```
Epoch 1: Train err: 0.488375, Train loss: 0.6930624028993031 |Validation err: 0.488375
Epoch 2: Train err: 0.48625, Train loss: 0.6930225396913196 |Validation err: 0.48625
Epoch 3: Train err: 0.484875, Train loss: 0.6930104787387545 |Validation err: 0.484875
Epoch 4: Train err: 0.48375, Train loss: 0.6929737159184047 |Validation err: 0.48375
Epoch 5: Train err: 0.481625, Train loss: 0.6929437924945165 |Validation err: 0.481625
Epoch 6: Train err: 0.480375, Train loss: 0.6929201266122242 |Validation err: 0.480375
Epoch 7: Train err: 0.4775, Train loss: 0.6928898332610963 |Validation err: 0.490625
Epoch 8: Train err: 0.47625, Train loss: 0.6928518338808938 |Validation err: 0.47625
Epoch 9: Train err: 0.473375, Train loss: 0.6928283364053757 |Validation err: 0.473375
Epoch 10: Train err: 0.4735, Train loss: 0.6927830217376588 |Validation err: 0.4735
Epoch 11: Train err: 0.471375, Train loss: 0.692752634722089 |Validation err: 0.471375
Epoch 12: Train err: 0.47, Train loss: 0.6927369047725012 |Validation err: 0.4825
Epoch 13: Train err: 0.46775, Train loss: 0.6927011656382728 |Validation err: 0.46775
Epoch 14: Train err: 0.466, Train loss: 0.6926825481747824 |Validation err: 0.47625
Epoch 15: Train err: 0.465375, Train loss: 0.6926437039223928 |Validation err: 0.465375
Epoch 16: Train err: 0.465875, Train loss: 0.6926113017021663 |Validation err: 0.465875
Epoch 17: Train err: 0.4655, Train loss: 0.6925884409556313 |Validation err: 0.4655
Epoch 18: Train err: 0.464, Train loss: 0.6925563386508397 |Validation err: 0.464
Epoch 19: Train err: 0.462625, Train loss: 0.6925268693575783 |Validation err: 0.462625
Epoch 20: Train err: 0.461, Train loss: 0.6924817524259053 |Validation err: 0.461
Epoch 21: Train err: 0.45875, Train loss: 0.6924625351315453 |Validation err: 0.45875
Epoch 22: Train err: 0.457375, Train loss: 0.6924205725155179 |Validation err: 0.457375
Epoch 23: Train err: 0.456625, Train loss: 0.6924104567558046 |Validation err: 0.456625
Epoch 24: Train err: 0.458125, Train loss: 0.692374030749003 |Validation err: 0.458125
Epoch 25: Train err: 0.4575, Train loss: 0.6923392566423567 |Validation err: 0.4575
Epoch 26: Train err: 0.455, Train loss: 0.6923072669241164 |Validation err: 0.4475
Epoch 27: Train err: 0.450875, Train loss: 0.6922921983022539 |Validation err: 0.450875
Epoch 28: Train err: 0.451875, Train loss: 0.6922409875052316 |Validation err: 0.451875
Epoch 29: Train err: 0.45225, Train loss: 0.6922226188674806 |Validation err: 0.45225
Epoch 30: Train err: 0.451625, Train loss: 0.6921830574671427 |Validation err: 0.451625
Finished Training
```



▼ Part 4. Evaluating the Best Model [15 pt]



▼ Part (a) - 1 pt

Choose the **best** model that you have so far. This means choosing the best model checkpoint, including the choice of `small_net` vs `large_net`, the `batch_size`, `learning_rate`, **and the epoch number**.

Modify the code below to load your chosen set of weights to the model object `net`.

```
net = LargeNet()
model_path = get_model_name("large", batch_size=128, learning_rate=0.005, epoch=29)
state = torch.load(model_path)
net.load_state_dict(state)
```

```
↳ <All keys matched successfully>
```

▼ Part (b) - 2pt

Justify your choice of model from part (a).

The large net seemed to perform just slightly better in terms of validation errors, although both nets were pretty similar. Aside from this, the test with the moderate batch size of 128 and learning rate of 0.005 produced good results so these were the parameters I chose.

▼ Part (c) - 2pt

Using the code in Part 0, any code from lecture notes, or any code that you write, compute and report the **test classification error** for your chosen model.

```
# If you use the `evaluate` function provided in part 0, you will need to
# set batch_size > 1
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=128)
```

```
testing = nn.BCEWithLogitsLoss()
```

```
test_err, test_loss = evaluate(net, test_loader, testing)
print(test_err, test_loss)
```

```
val_err, val_loss = evaluate(net, val_loader, testing)
print(val_err, val_loss)
```

```
↳ Files already downloaded and verified
Files already downloaded and verified
0.312 0.5874807238578796
0.3275 0.6071126200258732
```

The test classification error for the model is 0.312

▼ Part (d) - 3pt

How does the test classification error compare with the **validation error**? Explain why you would expect the test error to be *higher* than the validation error.

The test error is slightly less than the validation error which contradicts our expectations of a higher validation error. We fit the hyperparameters to the training data so technically the validation error would be less since the model has been trained to take on this data which hasn't been seen before.

▼ Part (e) - 2pt

Why did we only use the test data set at the very end? Why is it important that we use the test data as little as possible?

When a model is being trained it is very important that the test data is only used when you are actually ready to test the model. You want to feed your model fresh unseen data to see how it really performs in an unbiased way. You don't want to risk over/under fitting with the testing data because then you are just fitting the model to work with this data.

▼ Part (f) - 5pt

How does the your best CNN model compare with an 2-layer ANN model (no convolutional layers) on classifying cat and dog images. You can use a 2-layer ANN architecture similar to what you used in Lab 1. You should explore different hyperparameter settings to determine how well you can do on the validation dataset. Once satisfied with the performance, you may test it out on the test data.

Hint: The ANN in lab 1 was applied on greyscale images. The cat and dog images are colour (RGB) and so you will need to flattened and concatenate all three colour layers before feeding them into an ANN.

```
# Below is the ANN from Lab 1
```

```
import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import datasets, transforms
import matplotlib.pyplot as plt # for plotting
import torch.optim as optim
```

```
torch.manual_seed(1) # set the random seed
```

```
# define a 2-layer artificial neural network
class Pigeon(nn.Module):
    def __init__(self):
        super(Pigeon, self).__init__()
        self.layer1 = nn.Linear(3*32*32, 60)
        self.layer2 = nn.Linear(60,30)
        self.layer3 = nn.Linear(30,1)
        self.name = "pigeon"

    def forward(self, img):
        flattened = img.view(-1,3 * 32 * 32)
        activation1 = self.layer1(flattened)
        activation1 = F.relu(activation1)
        activation2 = self.layer2(activation1)
        activation2 = F.relu(activation2)
        activation3 = self.layer3(activation2)
        return activation3.squeeze(1)

pigeon = Pigeon()

train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=128)

train_net(pigeon,batch_size=128,learning_rate=0.005,num_epochs=30)

testing = nn.BCEWithLogitsLoss()

test_err, test_loss = evaluate(net, test_loader, testing)
print(test_err)
```

